

# Kraken 2

Improved metagenomic analysis with Kraken 2  
Derrick E. Wood, Jennifer Lu & Ben Langmead  
Genome Biol 20, 257 (2019)

<https://doi.org/10.1186/s13059-019-1891-0>

Konstantin Schütze  
Advanced Bioinformatics 1  
2021-09-15


# Metagenomic Sequence Classification

# Reads + Reference Genomes → Taxonomic assignment



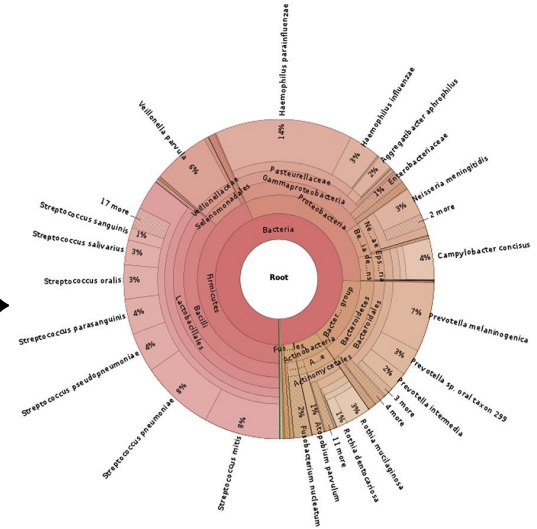
Illumina HiSeq 4000

+



A taxonomy

<https://doi.org/10.1038/s41598-018-30297-w>




Taxonomic distribution of saliva microbiome (kraken 1)

# Why do we even need that?


- Wild rat microbiome improves lab rat health  
<https://doi.org/10.1016/j.cell.2017.09.016>
- Zika virus evolution  
<https://doi.org/10.1038/nature22402>
- Finding the source of *Pseudomonas aeruginosa* infections in a hospital  
<http://dx.doi.org/10.1136/bmjopen-2014-006278>

# Rat Microbiome Diversity



High diversity in wild rats

Low diversity in lab rats



# Methods

1. **BLAST**: Search entire read
  - “Classic”
  - accurate
  - slow
2. **k-mers**: Split read into k-mers
  - faster
  - near-BLAST accuracy
3. **marker genes**: DB with only few genes (16S)
  - very fast
  - less accurate

# Kraken 1

1. Build fast mapping reference genomes k-mers -> taxonomy id
  - Minimizer
  - Offset table
2. Query all k-mers from read
3. Assign read by majority assignment of k-mers

# Minimizer (l-mer)

Minimizer: Lexicographically smallest l-mer inside the k-mer

ATGATCGTGCATCCATCAGTCATCAGTA <- initial k-mer

ATGATCGTGCATCCATCAGTCA

TGATCGTGCATCCATCAGTCAT

TATCGTGCATCCATCAGTCATC

ATCGTGCATCCATCAGTCATCA <- minimizer<sup>1</sup> (l-mer)

TCGTGCATCCATCAGTCATCAG

CGTGCATCCATCAGTCATCAGT

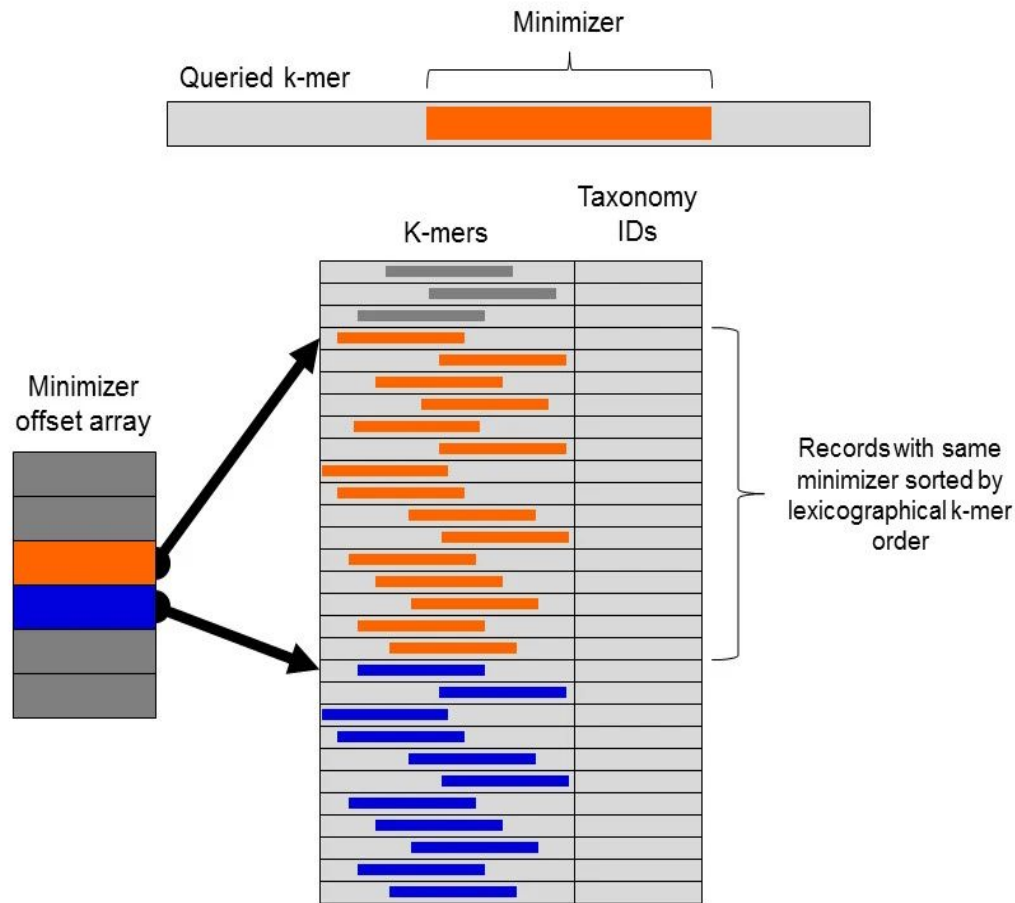
GTGCATCCATCAGTCATCAGTA

<sup>1</sup>ignoring XOR low complexity masking



# Kraken 1

- Group k-mers by minimizers
- Minimizer offset array
  - $8 \times 4^l$  bytes for l-mers
- Good for CPU caching



Can we compress k-mers even more?

# Spaced k-mers

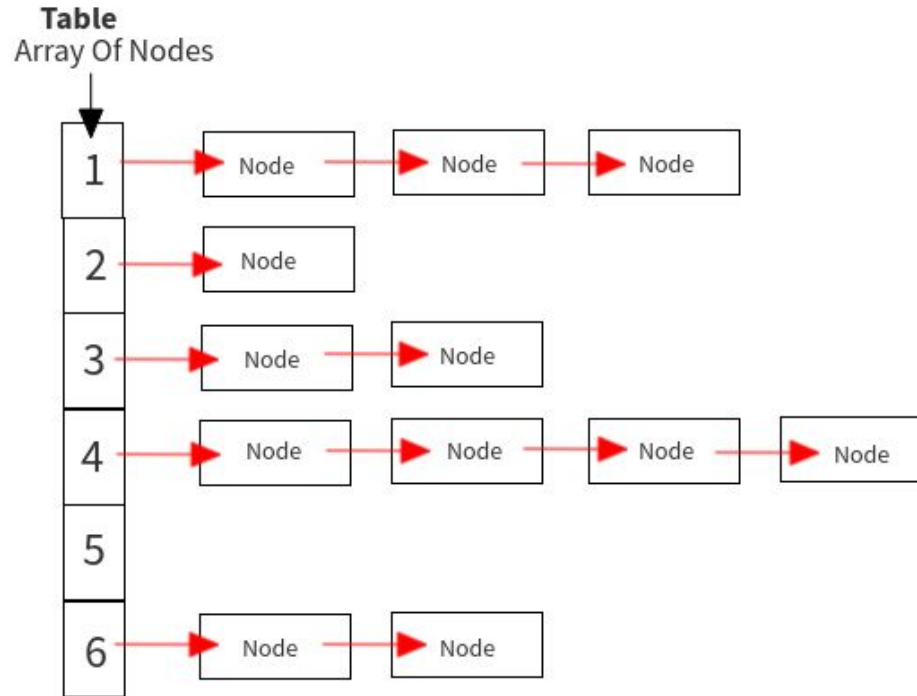
Mask  $s$  positions in  $l$ -mers (default:  $s=7$ )

AGTATCGTGCATCGATCAGTCATCA	AGTA	<- original k-mer
TCGTGCATCGATCAGTCATCA		<- initial minimizer
TCGTGCATCGATCAG-C-T-A		<- spaced minimizer


# Kraken 2

- Spaced k-mers
- k-mers share minimizers -> store only distinct minimizers
  - all k-mers -> few l-mers
- Only store hash of minimizer
- Defaults: k=35 l=31

# Textbook Hash Maps



# Compact Hash Table and Linear Probing



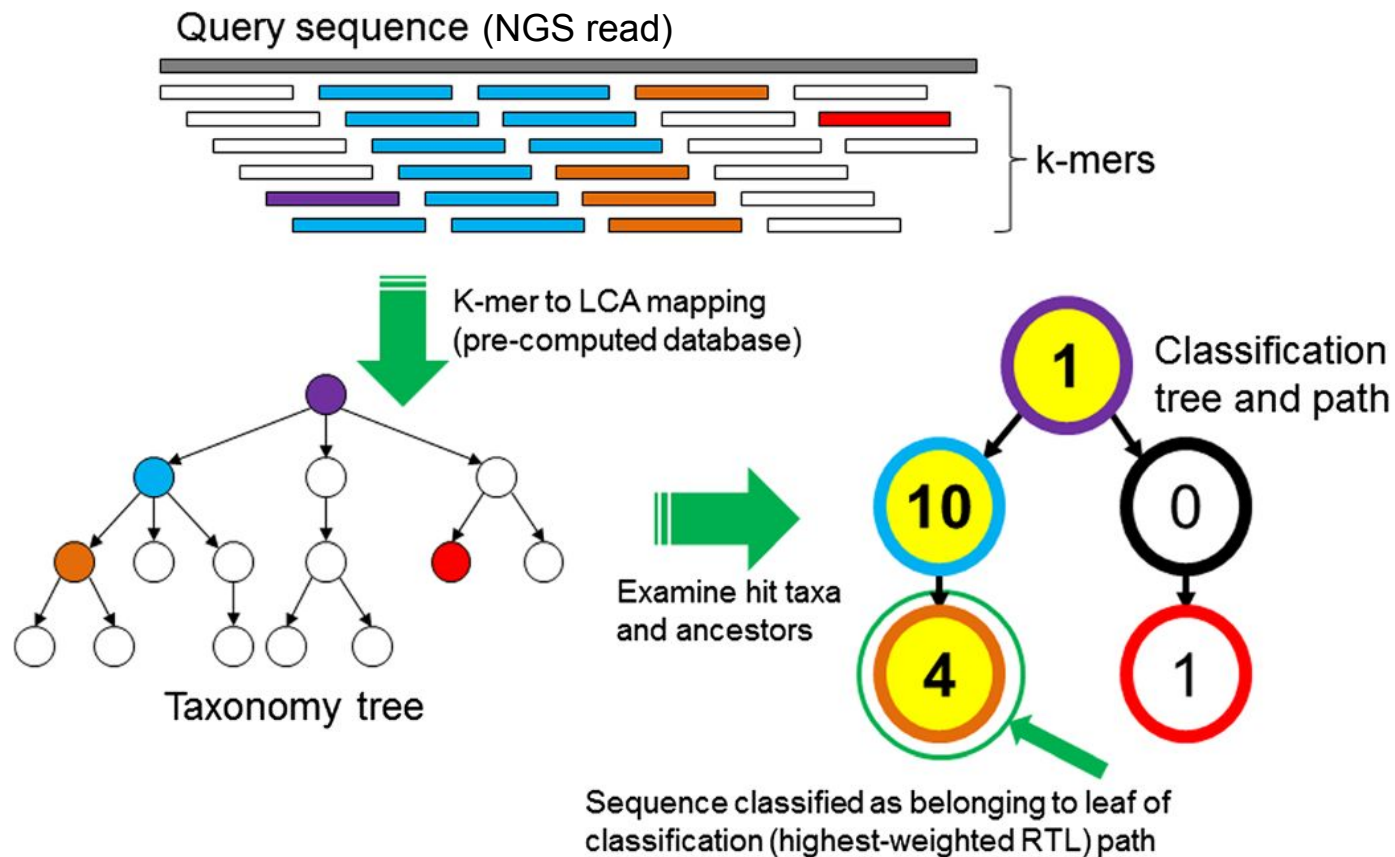
## Kraken 2: Memory Usage

- 4 bytes per entry: 15 bits hash, 17 bytes taxonomy id
- 70% load factor and D distinct minifier: **4D/0.7** bytes
- 85% less memory and 5x faster than Kraken 1





k-mer classification -> read classification



# Evaluation

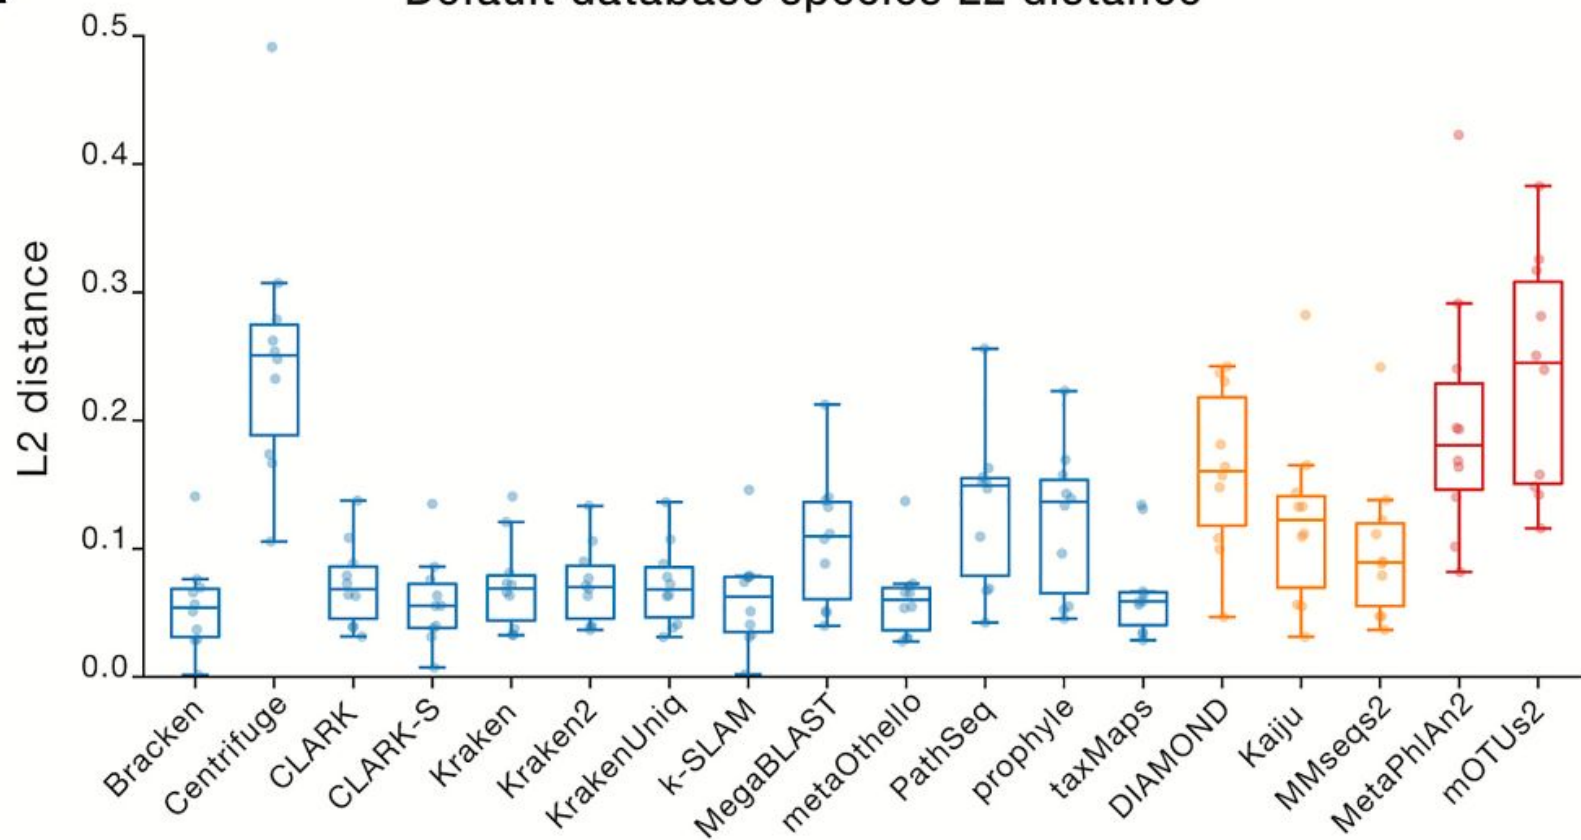
Type	Classifier	Memory Required	Time Required	Reference
DNA	CLARK-S	170 Gb	40 min	<a href="#">Ounit and Lonardi, 2016</a>
	Kraken	190 Gb	1 min	<a href="#">Wood and Salzberg, 2014</a>
	Kraken2	36 Gb	1 min	<a href="#">Wood and Salzberg, 2014</a>
	KrakenUniq	200 Gb	1 min	<a href="#">Breitwieser et al., 2018</a>
Protein	DIAMOND	110 Gb (varies)	10 min	<a href="#">Buchfink et al., 2015</a>
	MMseqs2	85 Gb (varies)	9 h	<a href="#">Steinegger and Söding, 2017</a>
Markers	MetaPhlAn2	2 Gb	1 min	<a href="#">Truong et al., 2015</a>


Benchmarking Metagenomics Tools for Taxonomic Classification, Ye et al. 2019

<https://doi.org/10.1016/j.cell.2019.07.010>

**A**

# Default database species L2 distance





High similarity between Kraken 1 and Kraken 2

# System Requirements and Colab

- 100 GB of disk space
- 29 GB of RAM
- “strongly suggest against using NFS storage”
- Use viral genomes instead

[https://colab.research.google.com/drive/1DHLIjov7hSTg1fliDOCErUvI2dW-EM6X?  
usp=sharing](https://colab.research.google.com/drive/1DHLIjov7hSTg1fliDOCErUvI2dW-EM6X?usp=sharing)

~End~





## Kraken 2

$k = 35, \ell = 31$

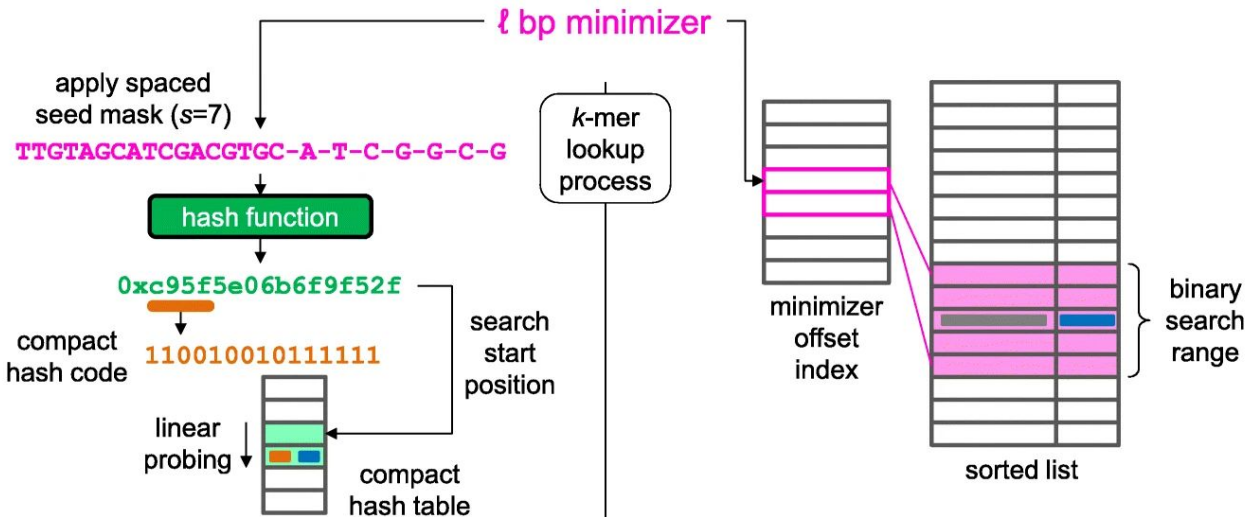
## Kraken 1

$k = 31, \ell = 15$


**a**

GCA**TTGTAGCATCGACGTGCCACTGCCGAGTCAGA**  $k$ -mer

**b**




**c**



# Hash Collisions? Not a Problem


	Load Factor								
	10%	20%	30%	40%	50%	60%	70%	80%	90%
26	0.000%	0.000%	0.000%	0.000%	0.000%	0.000%	0.000%	0.000%	0.000%
25	0.000%	0.000%	0.000%	0.000%	0.000%	0.000%	0.000%	0.000%	0.000%
24	0.000%	0.000%	0.000%	0.000%	0.000%	0.000%	0.000%	0.000%	0.000%
23	0.000%	0.000%	0.000%	0.000%	0.000%	0.000%	0.000%	0.000%	0.000%
22	0.000%	0.000%	0.000%	0.000%	0.000%	0.000%	0.000%	0.000%	0.001%
21	0.000%	0.000%	0.000%	0.000%	0.000%	0.000%	0.000%	0.001%	0.002%
20	0.000%	0.000%	0.000%	0.000%	0.000%	0.000%	0.000%	0.001%	0.005%
19	0.000%	0.000%	0.000%	0.000%	0.000%	0.001%	0.001%	0.002%	0.008%
18	0.000%	0.000%	0.000%	0.000%	0.001%	0.001%	0.001%	0.005%	0.017%
17	0.000%	0.000%	0.000%	0.000%	0.001%	0.002%	0.004%	0.009%	0.037%
16	0.000%	0.000%	0.001%	0.001%	0.002%	0.004%	0.008%	0.017%	0.074%
15	0.000%	0.001%	0.002%	0.003%	0.005%	0.009%	0.016%	0.036%	0.150%
14	0.000%	0.002%	0.004%	0.006%	0.009%	0.017%	0.032%	0.073%	0.296%
13	0.001%	0.003%	0.007%	0.012%	0.019%	0.032%	0.063%	0.145%	0.587%
12	0.002%	0.006%	0.012%	0.022%	0.038%	0.067%	0.123%	0.286%	1.145%
11	0.005%	0.013%	0.025%	0.044%	0.074%	0.129%	0.245%	0.555%	2.207%
10	0.012%	0.026%	0.049%	0.086%	0.150%	0.260%	0.480%	1.101%	4.030%
9	0.022%	0.052%	0.100%	0.174%	0.297%	0.505%	0.941%	2.132%	6.860%
8	0.045%	0.105%	0.195%	0.345%	0.574%	0.992%	1.820%	3.966%	10.895%
7	0.091%	0.216%	0.397%	0.686%	1.124%	1.931%	3.448%	6.936%	15.927%
6	0.182%	0.430%	0.783%	1.341%	2.183%	3.647%	6.216%	11.358%	21.765%



# Why do we even need that?

- Antibiotics can harm stem cell transplantation recipients

<https://doi.org/10.1126/scitranslmed.aaf2311>



## Taxonomy



## Simple taxonomy with tax Ids

(b)

Insert 8-bit minimizer hash with 0011 as least-sig bits, 0101 as most-sig (01010011), and 4 as taxon of origin. Immediately find empty slot.

(c)

Insert **11000100**, with **5**  
as taxon of origin.  
Immediately find empty  
slot.

(d)

Insert 01010011, with 6 as taxon of origin.  
Immediately find slot with matching most-sig bits, set value =  $LCA(4, 6) = 2$ .

(e)

Insert 11000011, with 7 as taxon of origin. Scan forward 1 row before finding match for 1100. Note: **this is a hash table collision**. Set value =  $LCA(5, 7) = 1$ .

(f)

Insert 10001100, with 8 as taxon of origin.  
Immediately find empty slot.

(g)

Insert **11001101**, with **9** as  
taxon of origin.  
Immediately find empty  
slot.

## Queries

(h)

most-sig(hash)	
LCA	
0	0
0	0
0	0
0101	2
1100	1
0	0
0	0
0	0
0	0
0	0
0	0
0	0
0	0
1000	8
1100	9
0	0
0	0

Query 1100 points to row 9.

Query 8-bit minimizer hash 10001100. Find 8, the correct answer for that hash.

(i)

most-sig(hash)	
LCA	
0	0
0	0
0	0
0101	2
1100	1
0	0
0	0
0	0
0	0
0	0
0	0
0	0
0	0
1000	8
1100	9
0	0
0	0

Query 0011 points to row 2.

Query 01010011. Find 2, the LCA of the taxa for the two identical hashes we inserted in steps b and d.

(j)

most-sig(hash)	
LCA	
0	0
0	0
0	0
0101	2
1100	1
0	0
0	0
0	0
0	0
0	0
0	0
0	0
0	0
1000	8
1100	9
0	0
0	0

Query 0100 points to row 1.

Query 11000100. Find 1, the LCA of the taxa for two hashes, one matching the query and one (11000011) not. These were added in steps c and e.

(k)

most-sig(hash)	
LCA	
0	0
0	0
0	0
0101	2
1100	1
0	0
0	0
0	0
0	0
0	0
0	0
0	0
0	0
1000	8
1100	9
0	0
0	0

Query 0011 points to row 1.

Query 11000011. Scan forward 1 row before finding matching least-sig bits. Find 1, the LCA of the taxa for two hashes, one matching the query and one (11000100) not. These were added in steps c and e.

(l)

most-sig(hash)	
LCA	
0	0
0	0
0	0
0101	2
1100	1
0	0
0	0
0	0
0	0
0	0
0	0
0	0
0	0
1000	8
1100	9
0	0
0	0

Query 1100 points to row 9.

Query 11001100. Scan forward 1 row before finding matching least-sig bits. This yields the LCA 9, which is **spurious because this query hash was never inserted**.